

# 强化学习与博弈论

## Reinforcement Learning and Game Theory

陈旭

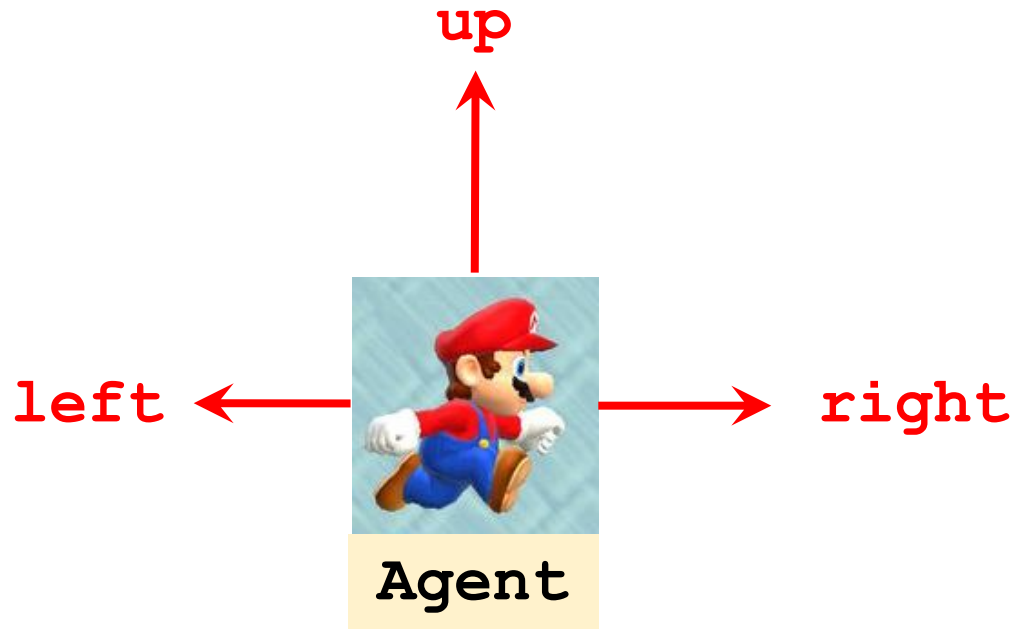
计算机学院



中山大學  
SUN YAT-SEN UNIVERSITY

# **Deterministic Policy Gradient RL**

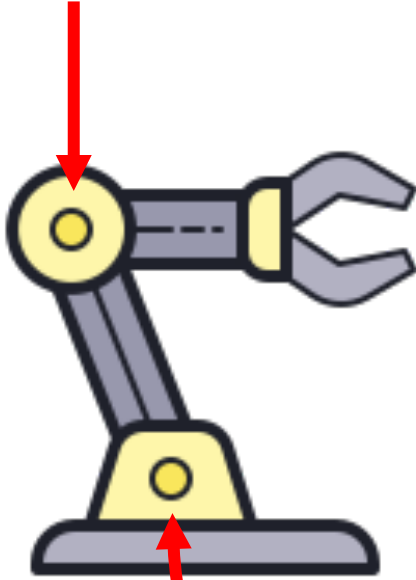
# Discrete Action Space



- Action space  $\mathcal{A} = \{\text{left}, \text{right}, \text{up}\}$ .
- The action space  $\mathcal{A}$  is discrete.

# Continuous Action Space

$$a_1 \in [0^\circ, 360^\circ]$$



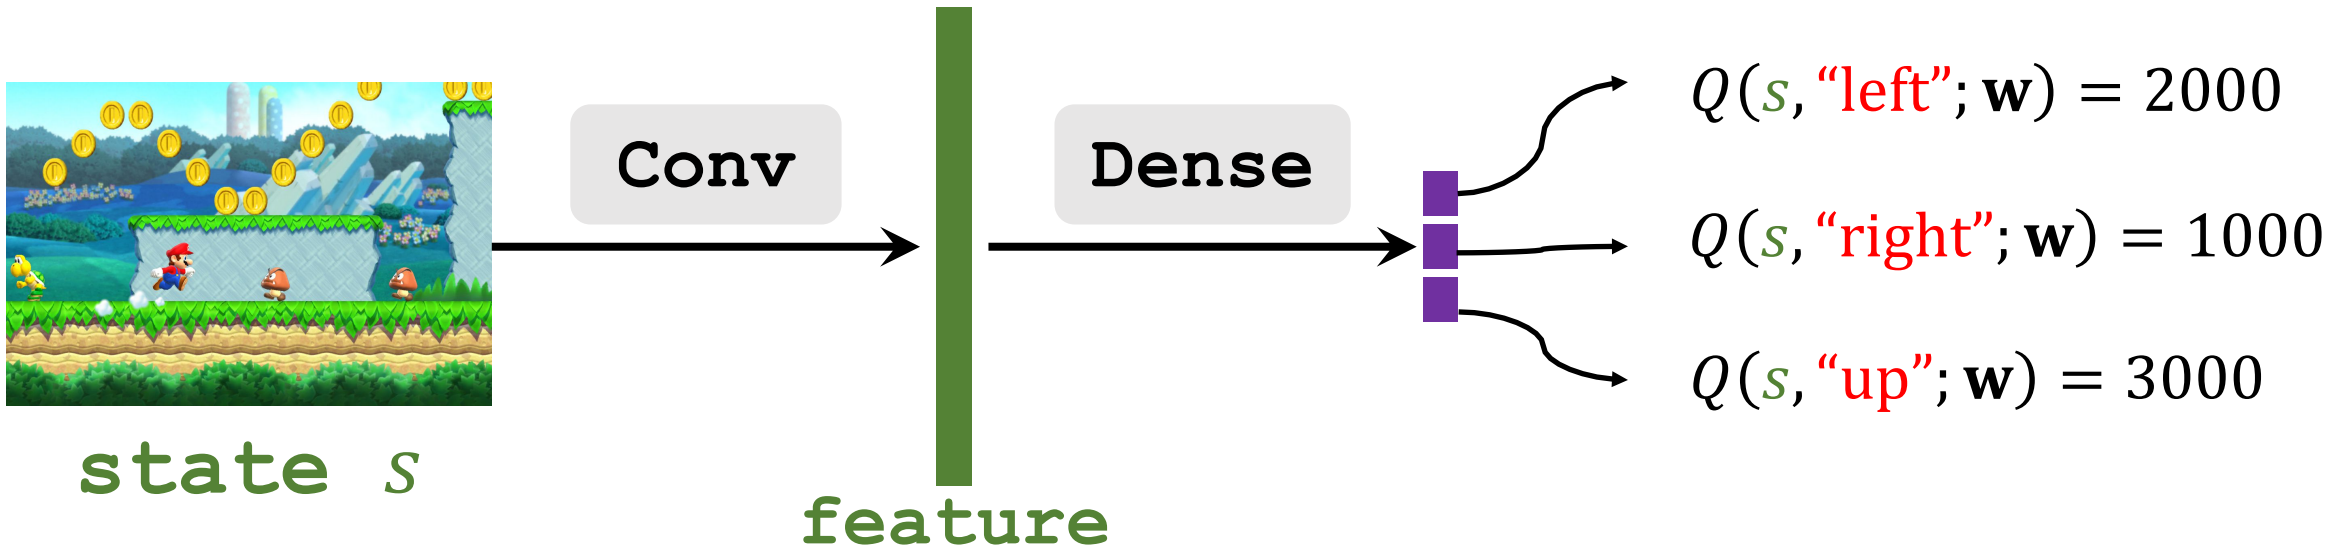
$$a_2 \in [0^\circ, 180^\circ]$$

- The action space  $\mathcal{A}$  is continuous:

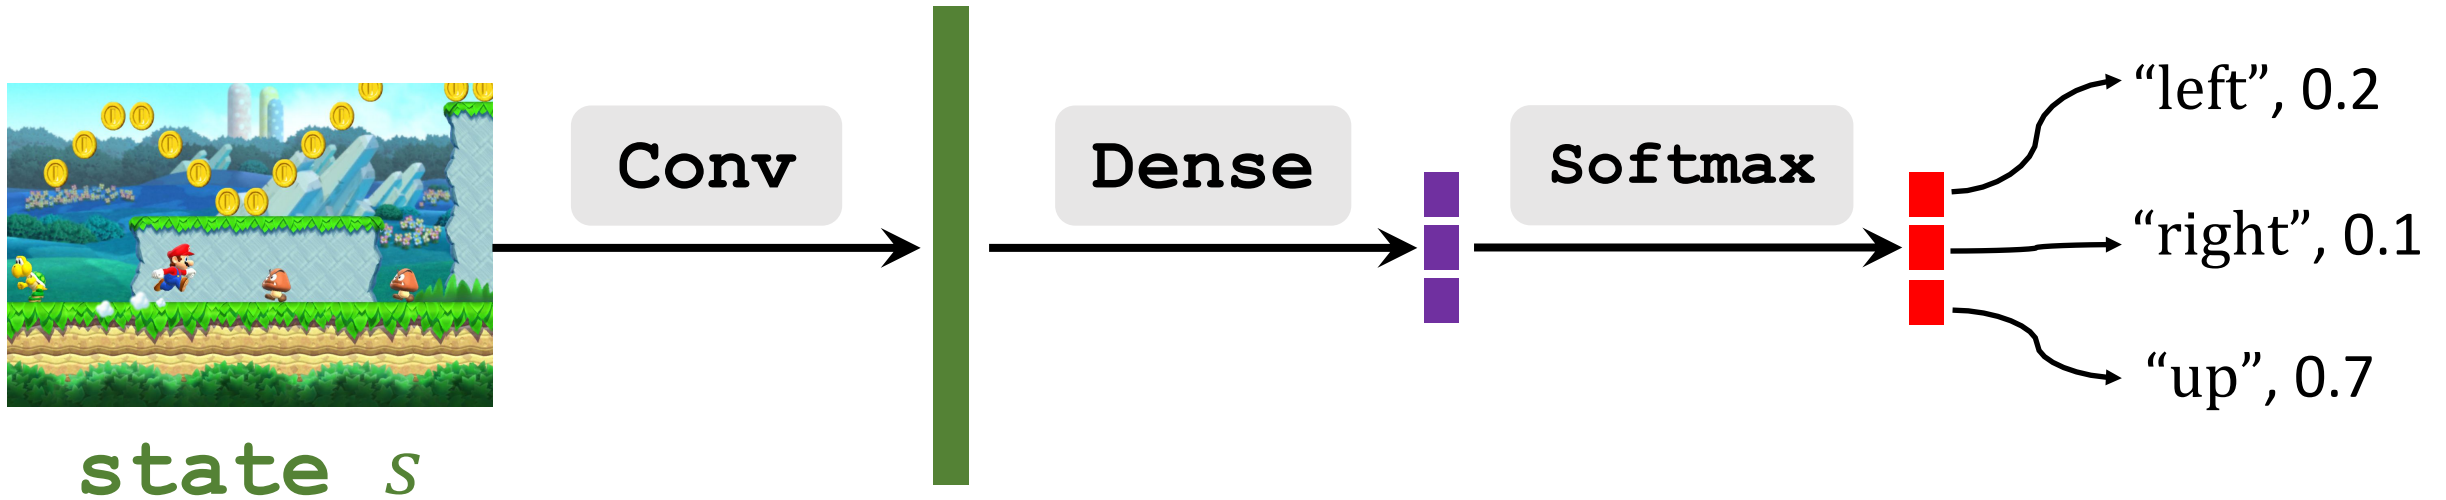
$$\mathcal{A} = [0^\circ, 360^\circ] \times [0^\circ, 180^\circ].$$

- Actions are 2-dim vectors.

# DQN for Discrete Action Space



# Policy Network for Discrete Action Space

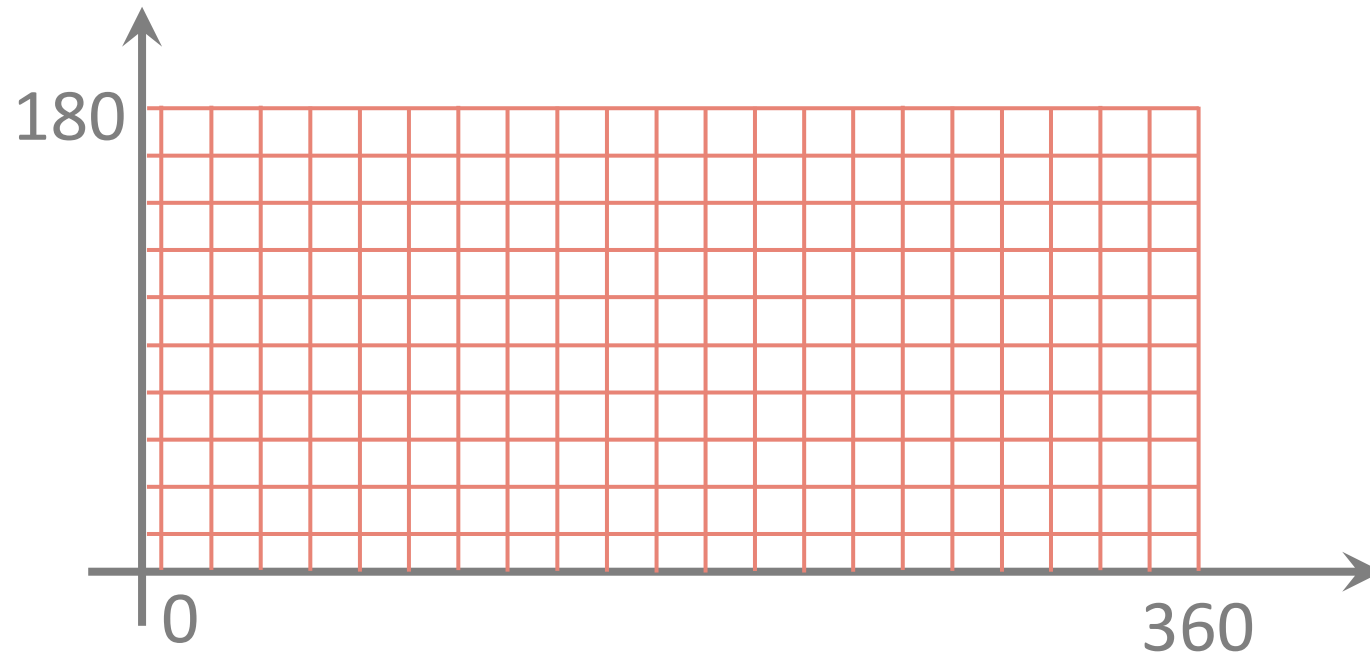


# Discretization



# Discretization

- Discretize the action space. (Draw a grid.)
- Now, the number of actions is the number of grid points.



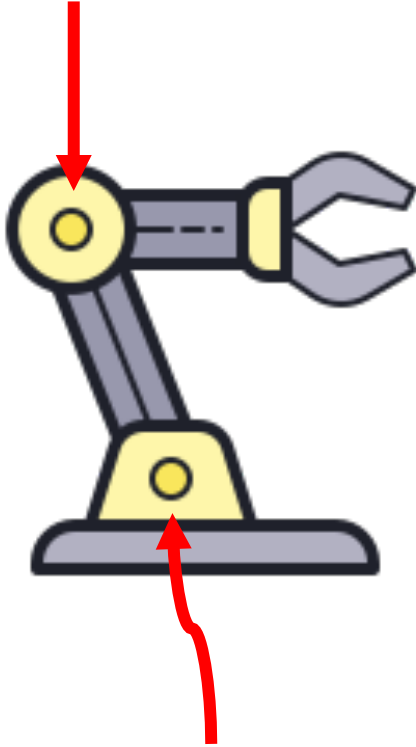


# Discretization

- Discretize the action space. (Draw a grid.)
- Now, the number of actions is the number of grid points.
- Problem: curse of dimensionality.
  - Let  $d$  be the degree of freedom.
  - The number of actions grows exponentially with  $d$ .

# Continuous Action Space

$$a_1 \in [0^\circ, 360^\circ]$$



$$a_2 \in [0^\circ, 180^\circ]$$

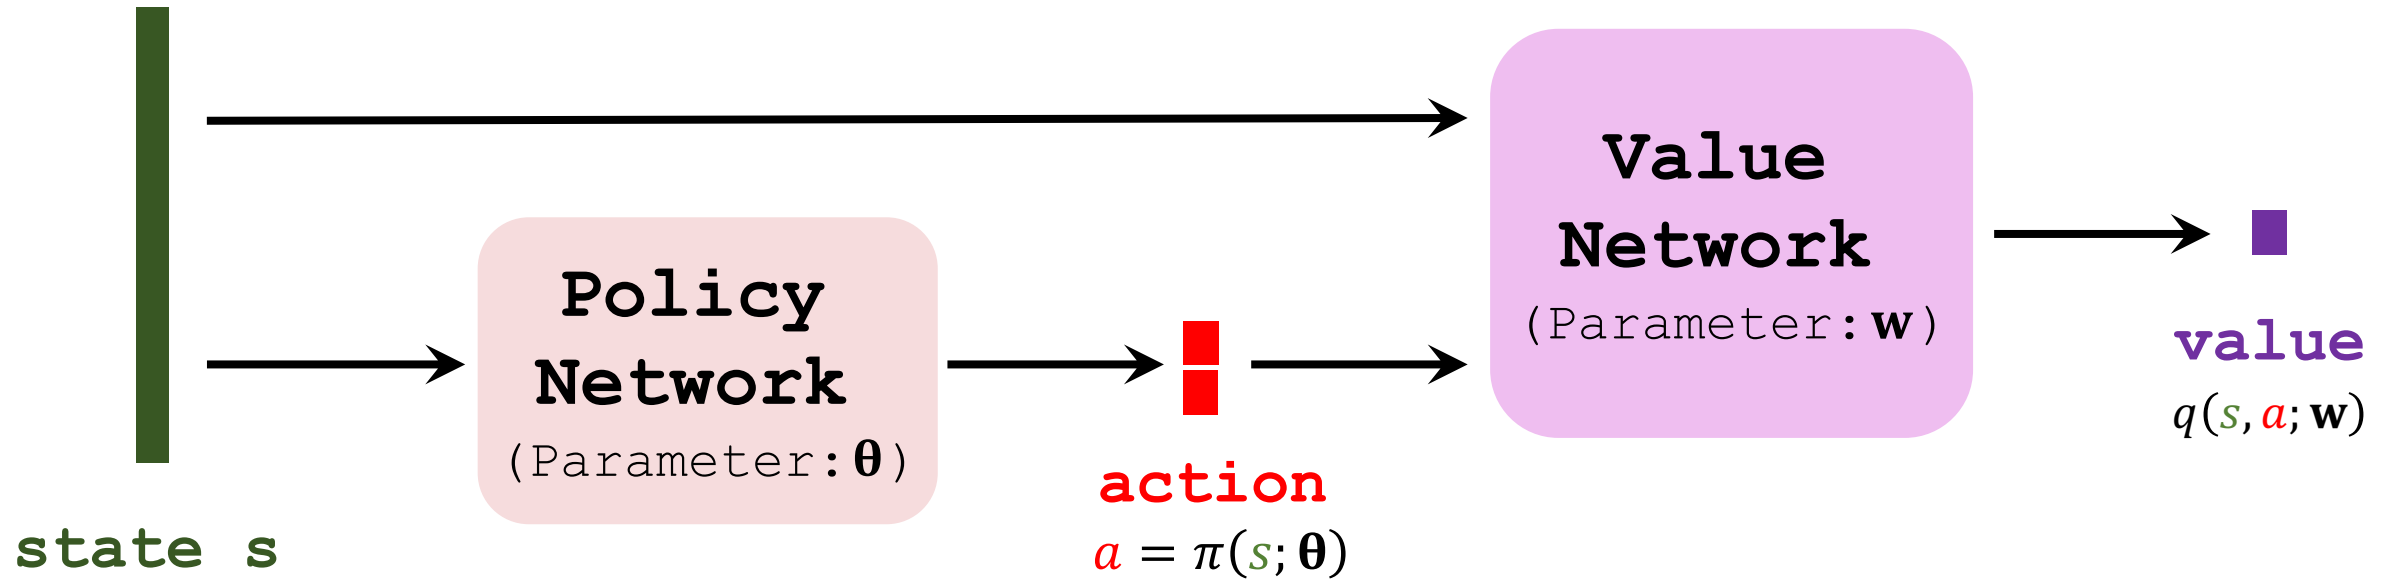
- The action space  $\mathcal{A}$  is a subset of  $\mathbb{R}^2$ .
- The action space  $\mathcal{A}$  is continuous:  
$$\mathcal{A} = [0^\circ, 360^\circ] \times [0^\circ, 180^\circ].$$
- Actions are 2-dim vectors.

# Deterministic Policy Gradient (DPG)

## Reference:

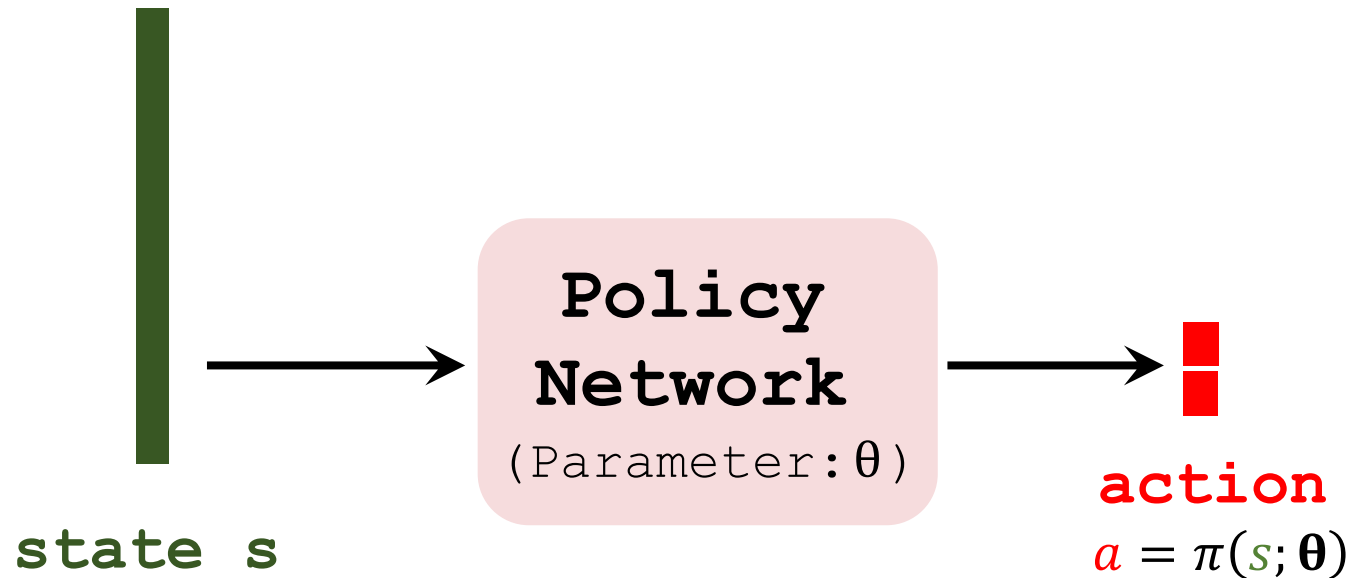
- Silver et al. [Deterministic policy gradient algorithms](#). In *ICML*, 2014.
- Lillicrap et al. [Continuous control with deep reinforcement learning](#). In *ICLR*, 2016.

# Deterministic Actor-Critic



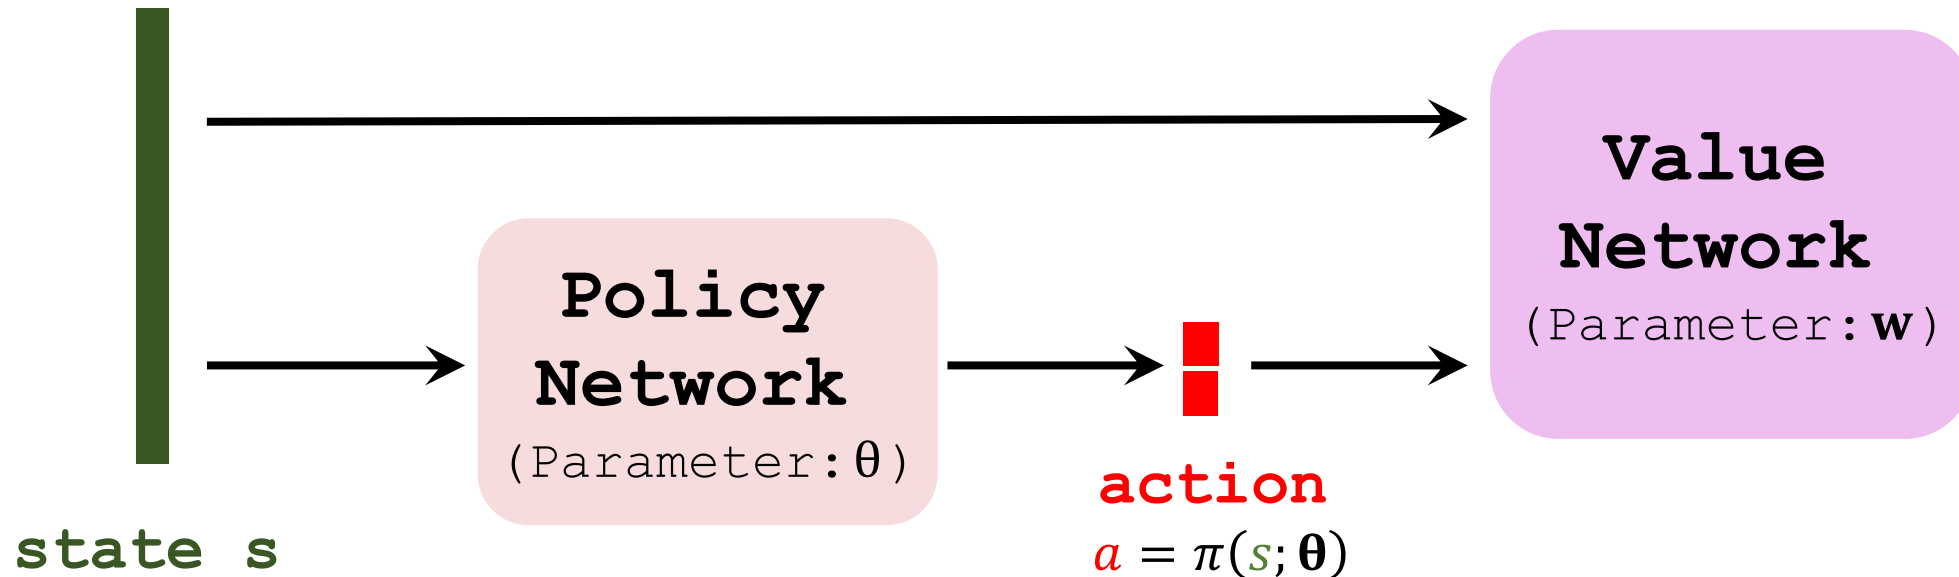
# Deterministic Actor-Critic

- Use a deterministic policy network (actor):  $a = \pi(s; \theta)$ .



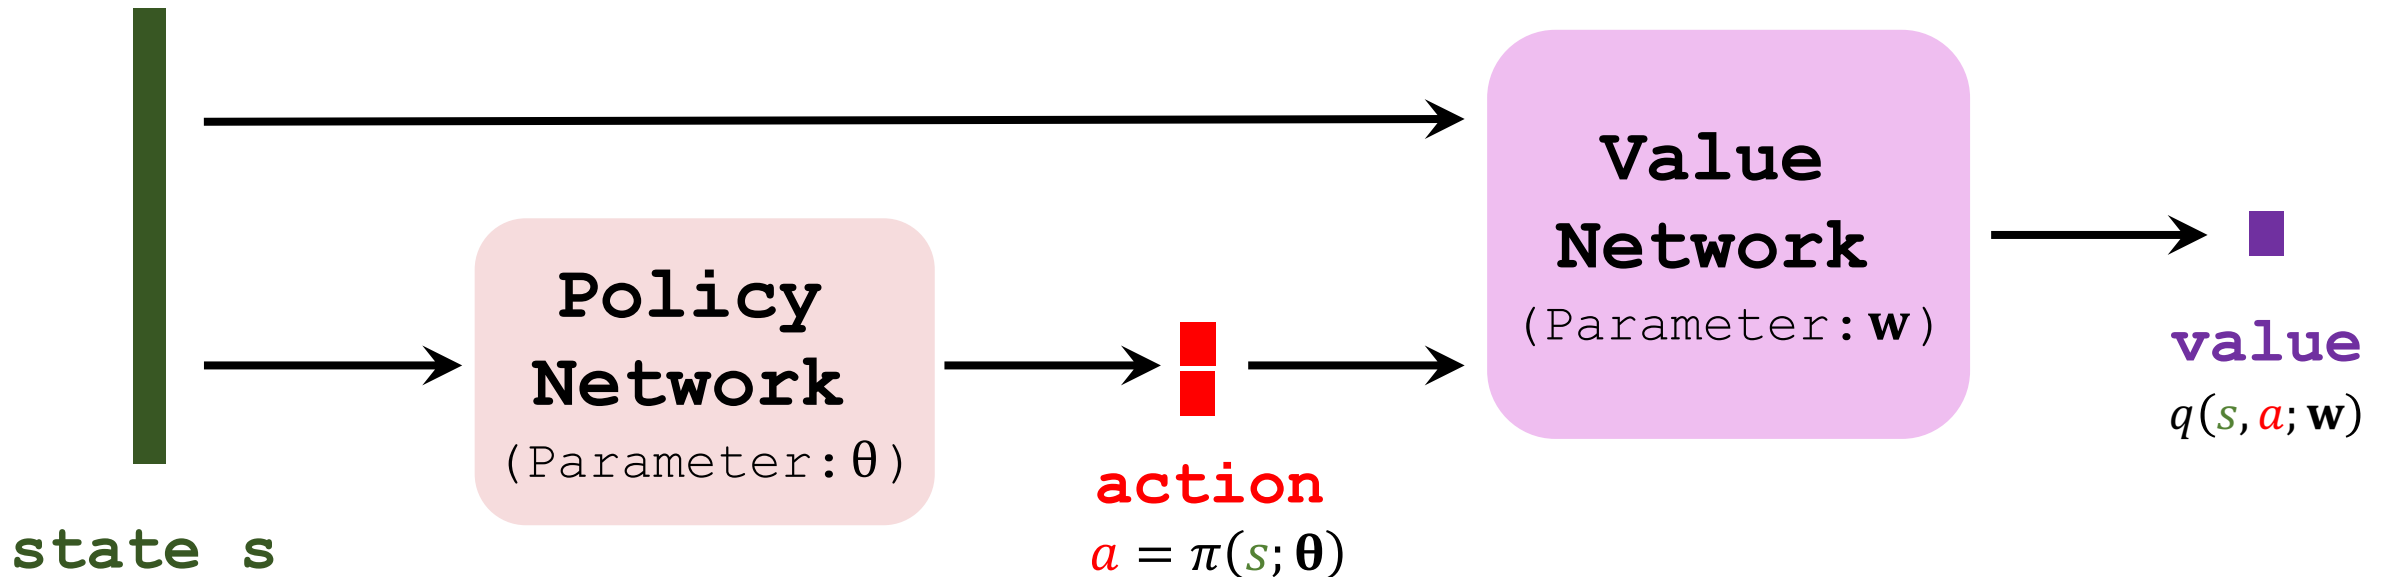
# Deterministic Actor-Critic

- Use a deterministic policy network (actor):  $a = \pi(s; \theta)$ .
- Use a value network (critic):  $q(s, a; \mathbf{w})$ .



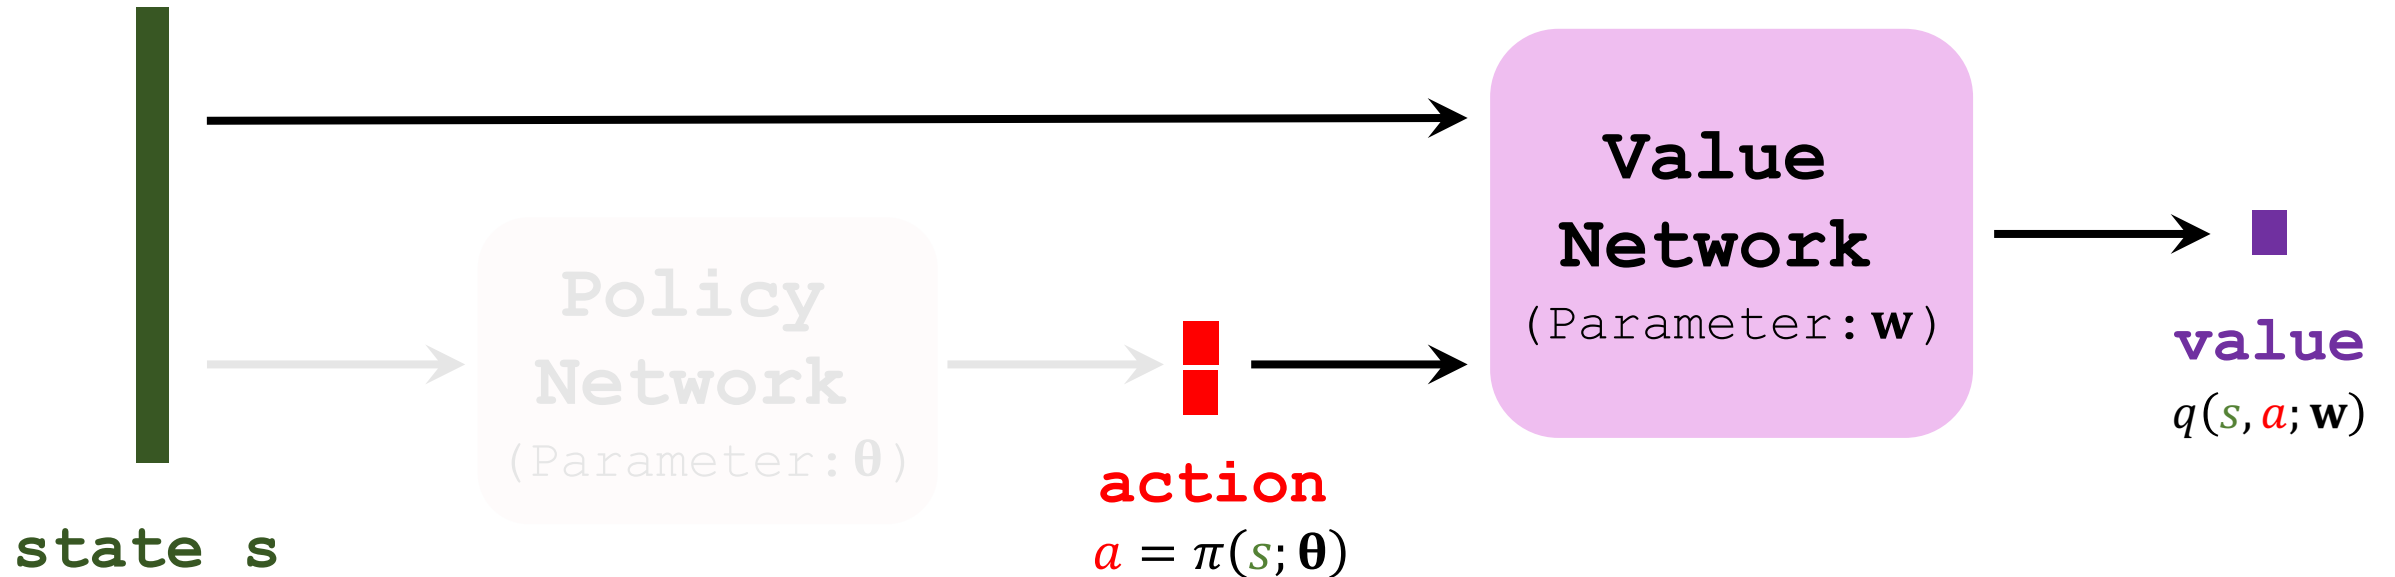
# Deterministic Actor-Critic

- Use a deterministic policy network (actor):  $a = \pi(s; \theta)$ .
- Use a value network (critic):  $q(s, a; w)$ .
- The critic outputs a scalar that evaluates *how good the action  $a$  is*.



# Updating Value Network by TD

- Transition:  $(s_t, a_t, r_t, s_{t+1})$ .





# Updating Value Network by TD

- Transition:  $(s_t, a_t, r_t, s_{t+1})$ .
- Value network makes prediction for time  $t$ :

$$q_t = q(s_t, a_t; \mathbf{w}).$$

# Updating Value Network by TD

- Transition:  $(s_t, a_t, r_t, s_{t+1})$ .
- Value network makes prediction for time  $t$ :

$$q_t = q(s_t, a_t; \mathbf{w}).$$

- Value network makes prediction for time  $t + 1$ :

$$q_{t+1} = q(s_{t+1}, a'_{t+1}; \mathbf{w}), \text{ where } a'_{t+1} = \pi(s_{t+1}; \boldsymbol{\theta}).$$

# Updating Value Network by TD

- Transition:  $(s_t, a_t, r_t, s_{t+1})$ .
- Value network makes prediction for time  $t$ :

$$q_t = q(s_t, a_t; \mathbf{w}).$$

- Value network makes prediction for time  $t + 1$ :

$$q_{t+1} = q(s_{t+1}, a'_{t+1}; \mathbf{w}), \text{ where } a'_{t+1} = \pi(s_{t+1}; \boldsymbol{\theta}).$$

- TD error:  $\delta_t = q_t - (r_t + \gamma \cdot q_{t+1})$ .



TD Target

# Updating Value Network by TD

- Transition:  $(s_t, a_t, r_t, s_{t+1})$ .
- Value network makes prediction for time  $t$ :

$$q_t = q(s_t, a_t; \mathbf{w}).$$

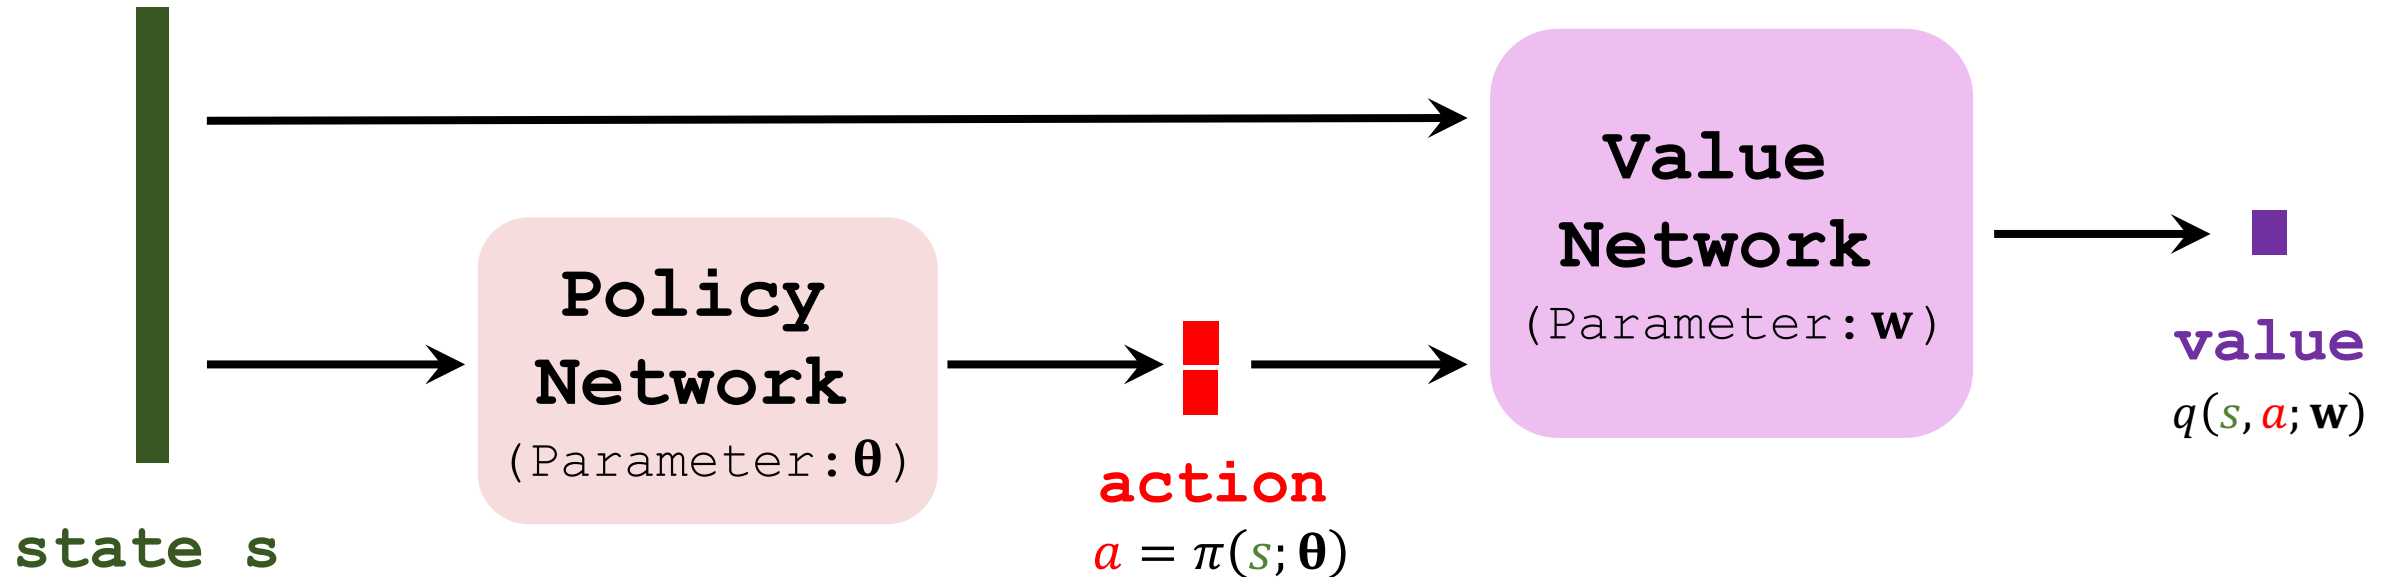
- Value network makes prediction for time  $t + 1$ :

$$q_{t+1} = q(s_{t+1}, a'_{t+1}; \mathbf{w}), \text{ where } a'_{t+1} = \pi(s_{t+1}; \boldsymbol{\theta}).$$

- TD error:  $\delta_t = q_t - (r_t + \gamma \cdot q_{t+1})$ .
- Update:  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \cdot \delta_t \cdot \frac{\partial q(s_t, a_t; \mathbf{w})}{\partial \mathbf{w}}$ .

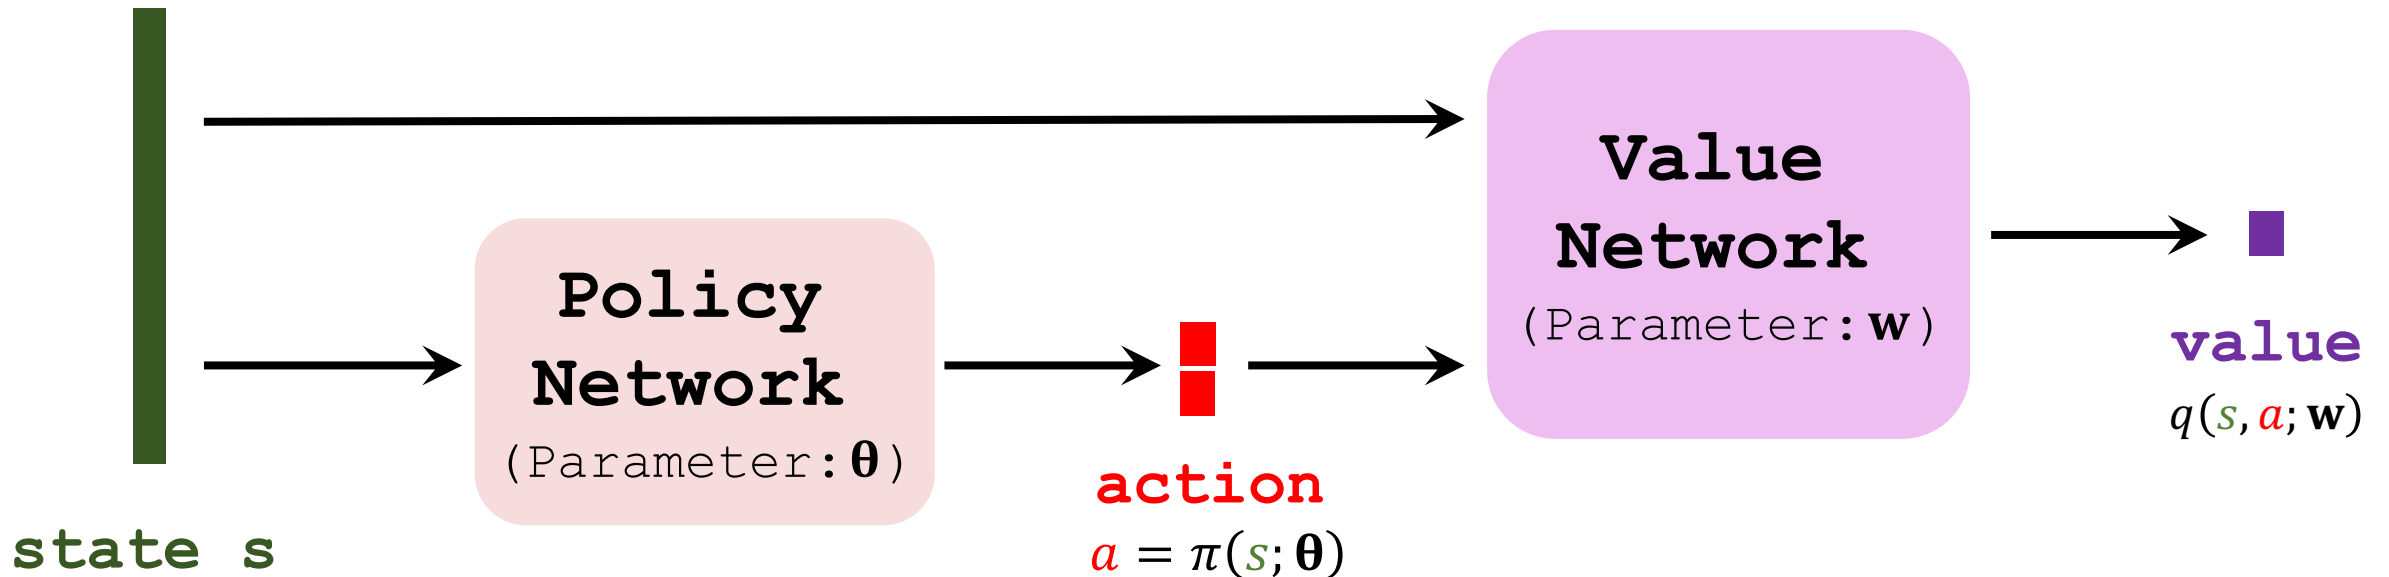
# Updating Policy Network by DPG

- The critic  $q(s, a; \mathbf{w})$  evaluates how good the action  $a$  is.



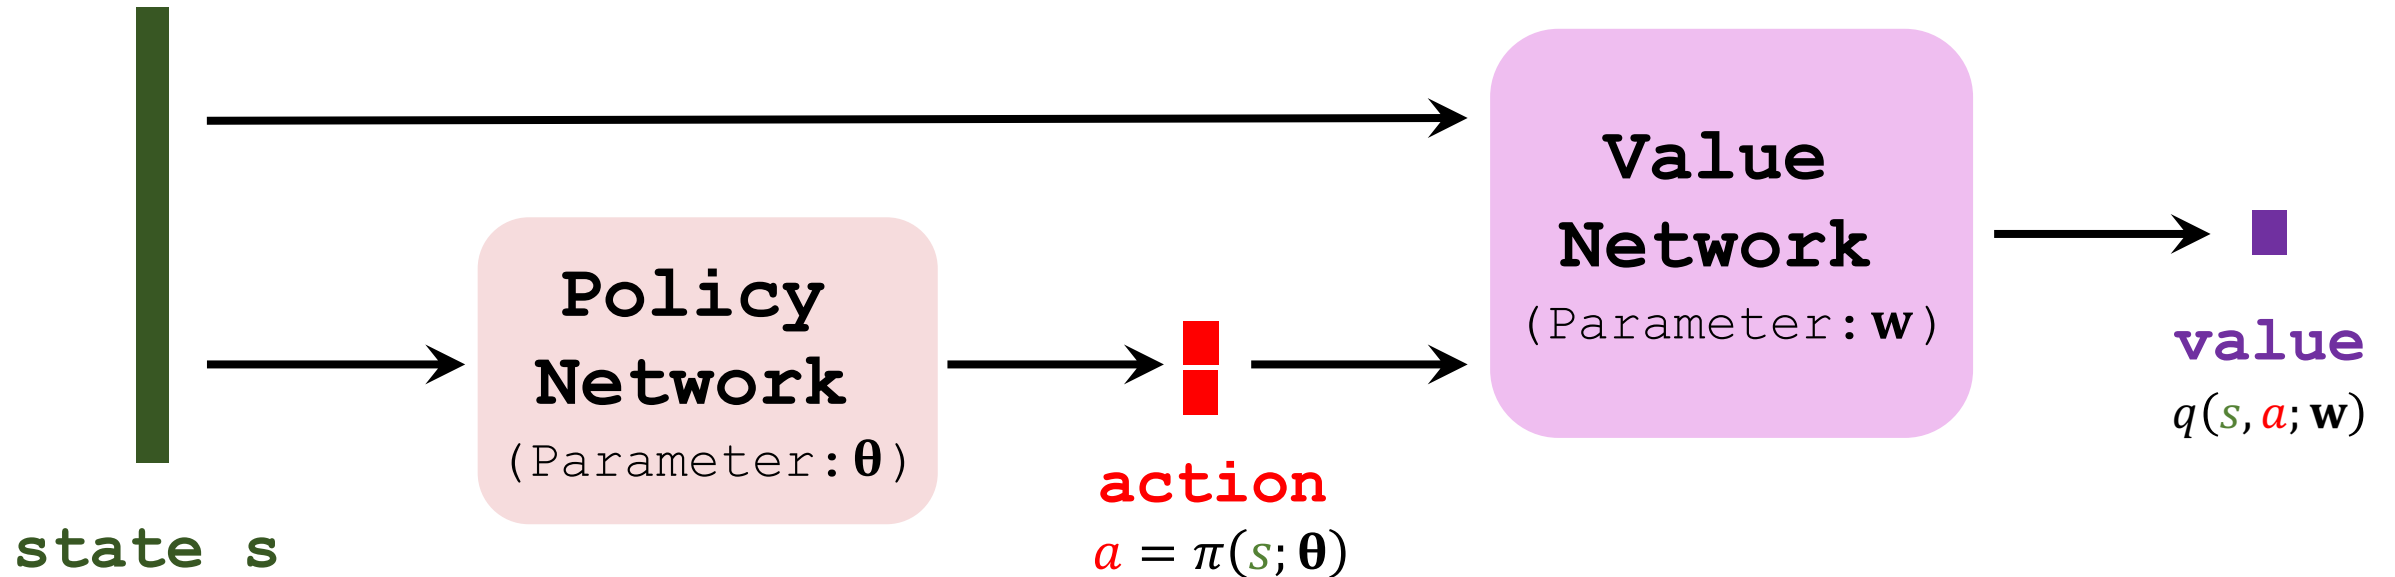
# Updating Policy Network by DPG

- The critic  $q(s, a; \mathbf{w})$  evaluates how good the action  $a$  is.
- Improve  $\theta$  so that the critic believes  $a = \pi(s; \theta)$  is better.
- Update  $\theta$  so that  $q(s, a; \mathbf{w}) = q(s, \pi(s; \theta); \mathbf{w})$  increases.



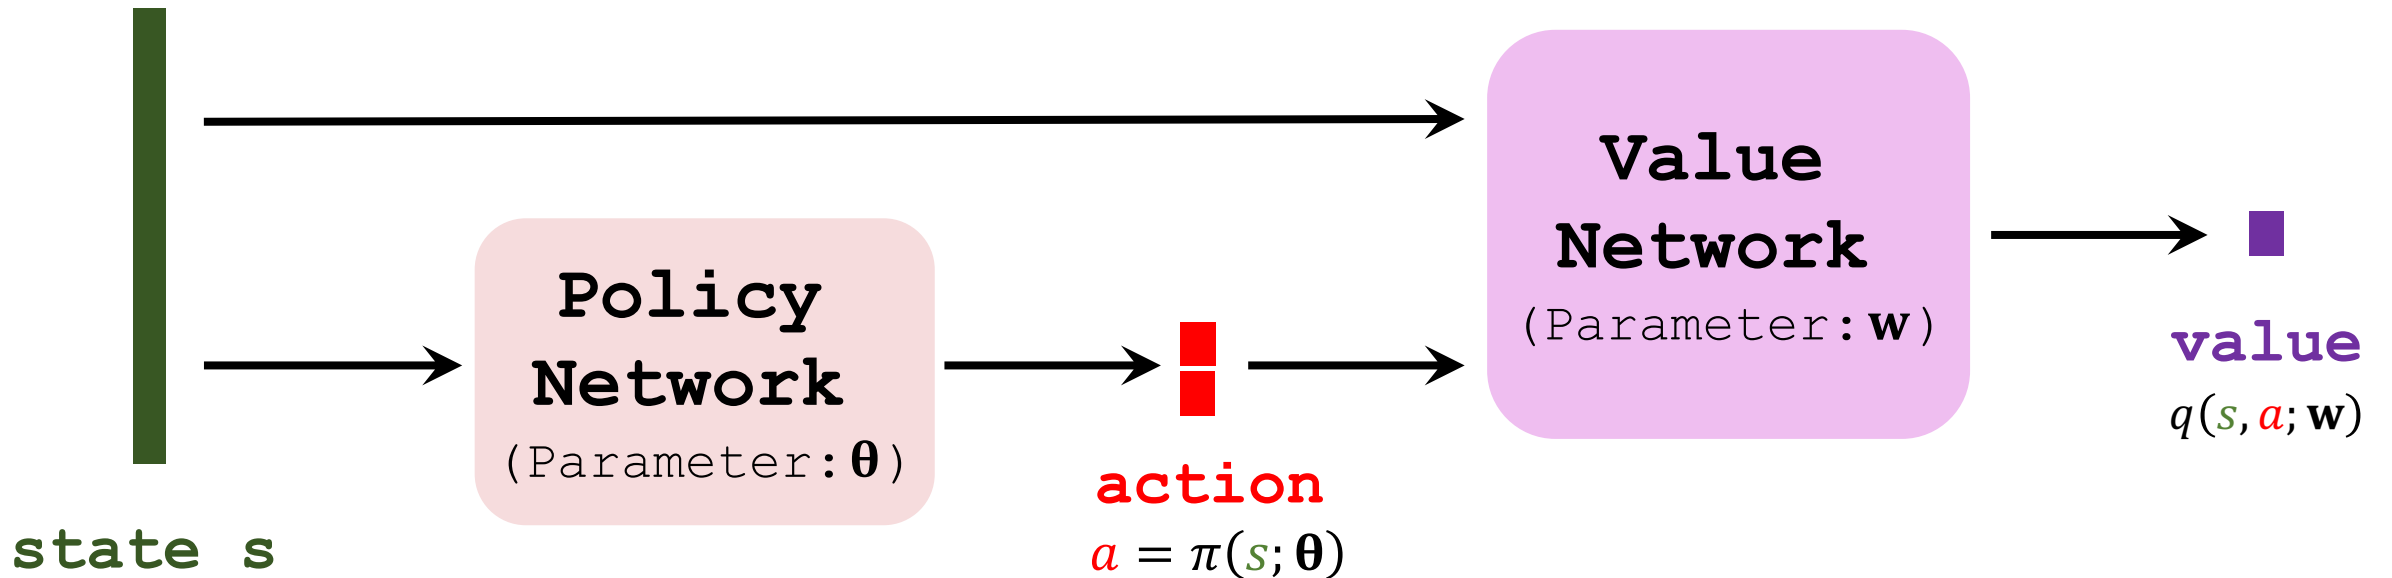
# Updating Policy Network by DPG

- **Goal:** Increasing  $q(s, a; \mathbf{w})$ , where  $a = \pi(s; \boldsymbol{\theta})$ .



# Updating Policy Network by DPG

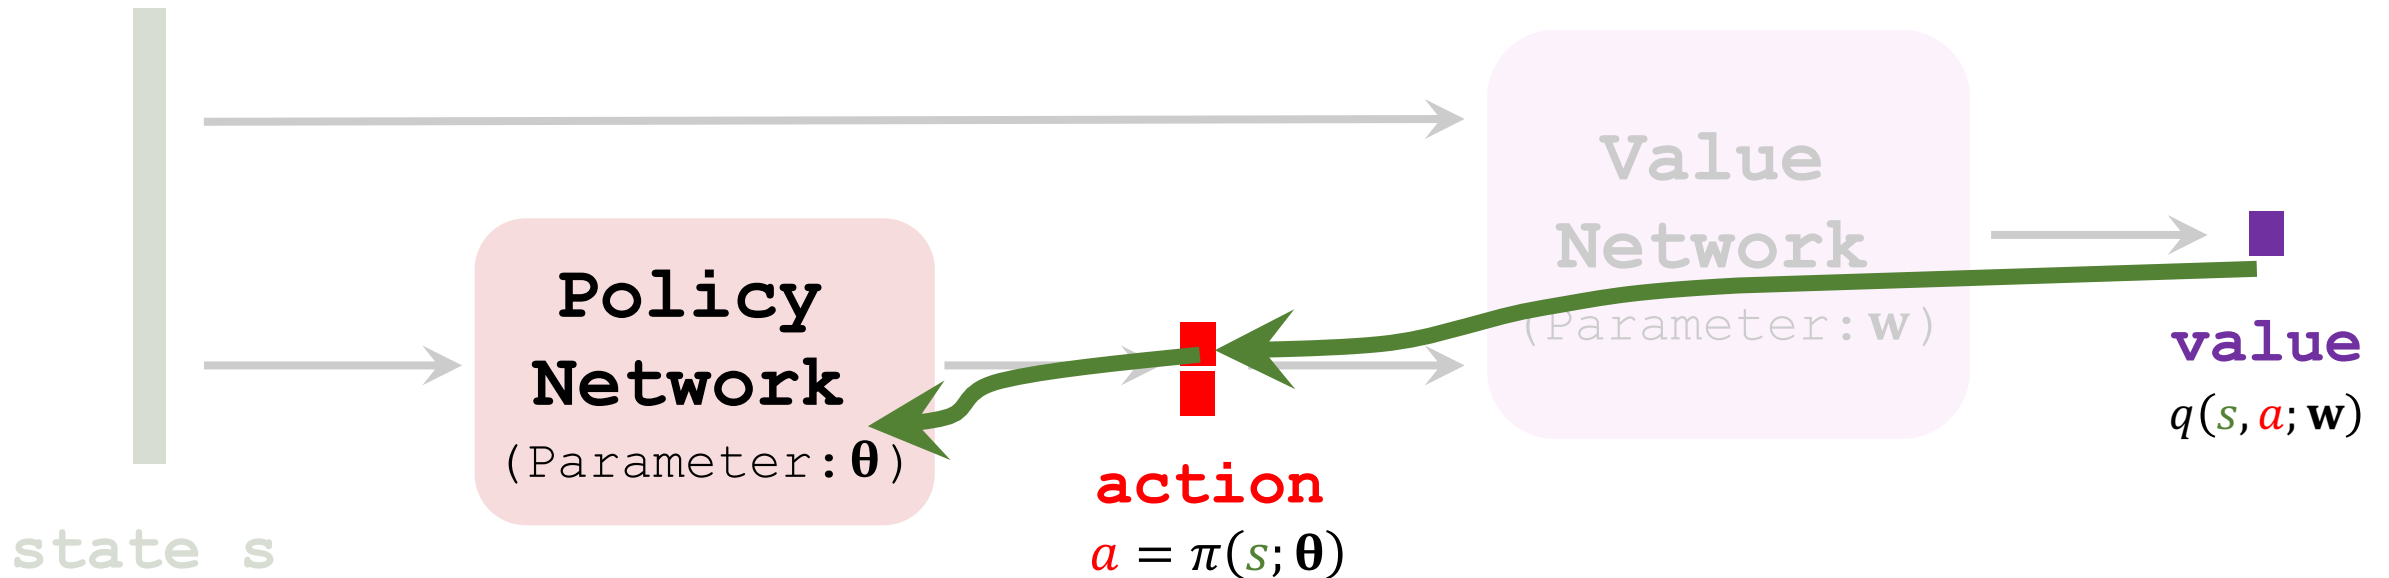
- **Goal:** Increasing  $q(s, a; \mathbf{w})$ , where  $a = \pi(s; \theta)$ .
- DPG:  $\mathbf{g} = \frac{\partial q(s, \pi(s; \theta); \mathbf{w})}{\partial \theta} = \frac{\partial a}{\partial \theta} \cdot \frac{\partial q(s, a; \mathbf{w})}{\partial a}$ .





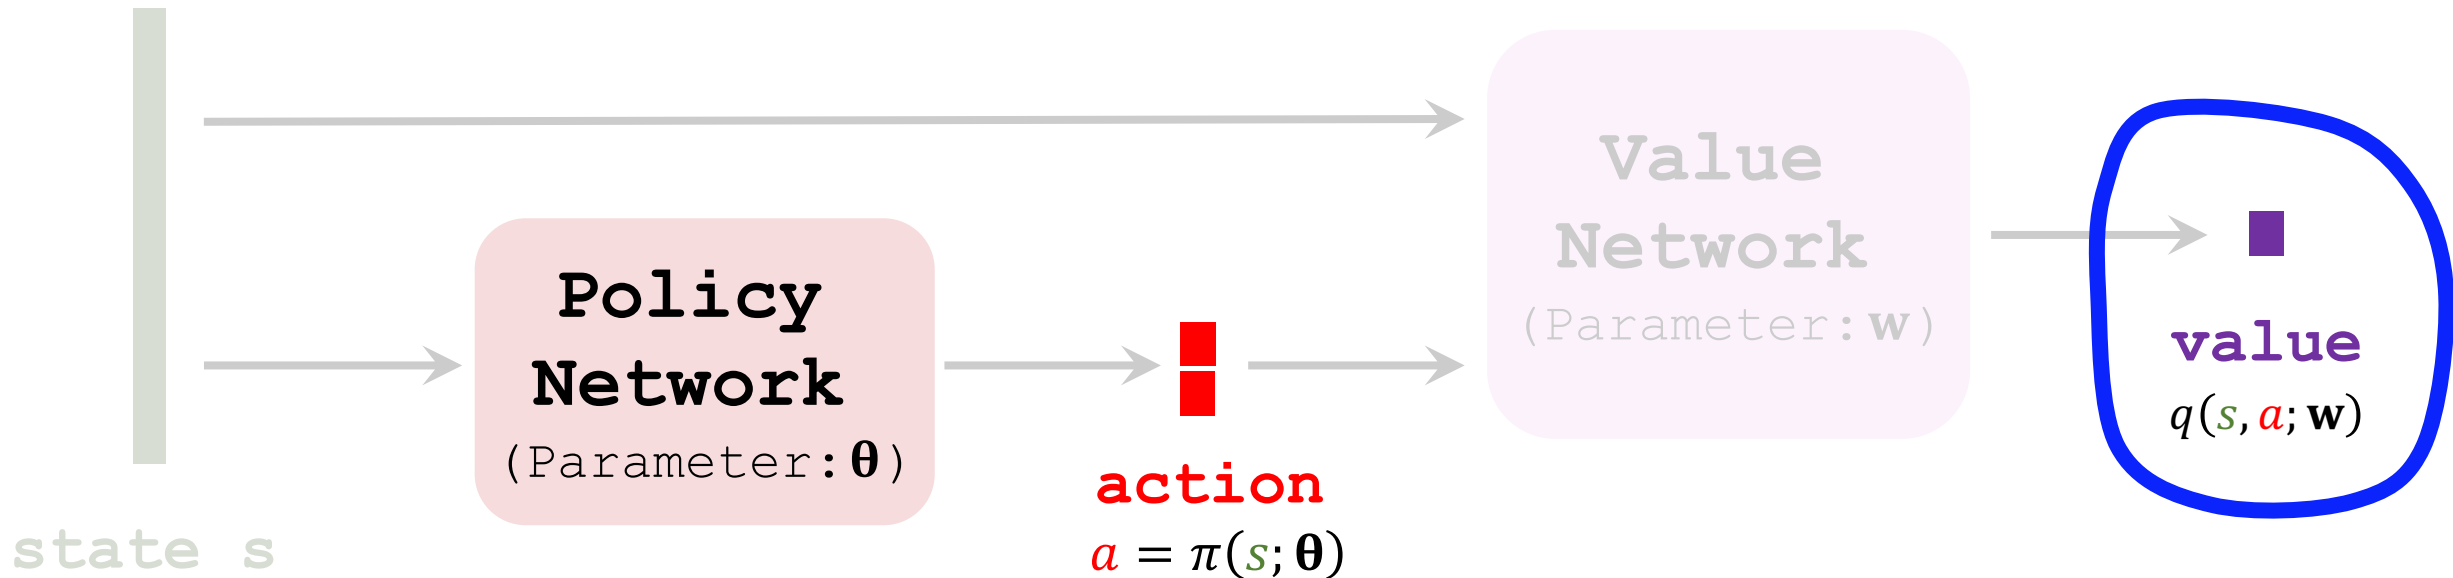
# Updating Policy Network by DPG

- **Goal:** Increasing  $q(s, a; \mathbf{w})$ , where  $a = \pi(s; \theta)$ .
- DPG:  $\mathbf{g} = \frac{\partial q(s, \pi(s; \theta); \mathbf{w})}{\partial \theta} = \frac{\partial a}{\partial \theta} \cdot \frac{\partial q(s, a; \mathbf{w})}{\partial a}$ .



# Updating Policy Network by DPG

- **Goal:** Increasing  $q(s, a; \mathbf{w})$ , where  $a = \pi(s; \theta)$ .
- DPG:  $\mathbf{g} = \frac{\partial q(s, \pi(s; \theta); \mathbf{w})}{\partial \theta} = \frac{\partial a}{\partial \theta} \cdot \frac{\partial q(s, a; \mathbf{w})}{\partial a}$ .
- Gradient ascent:  $\theta \leftarrow \theta + \beta \cdot \mathbf{g}$ .



# **Stochastic Policy VS Deterministic Policy**

## Stochastic Policy

## Deterministic Policy

Policy:

$$\pi(a|s; \theta)$$

$$\pi(s; \theta)$$

## Stochastic Policy

## Deterministic Policy

**Policy:**

$$\pi(\textcolor{red}{a}|\textcolor{green}{s}; \boldsymbol{\theta})$$

$$\pi(\textcolor{green}{s}; \boldsymbol{\theta})$$

**Output:**

Probability distribution  
over the action space

Action  $\textcolor{red}{a}$

## Stochastic Policy

## Deterministic Policy

**Policy:**

$$\pi(\textcolor{red}{a}|\textcolor{green}{s}; \boldsymbol{\theta})$$

$$\pi(\textcolor{green}{s}; \boldsymbol{\theta})$$

**Output:**

Probability distribution  
over the action space

Action  $\textcolor{red}{a}$

**Control:**

Randomly sample an action  
from the distribution

Directly use  
the output,  $\textcolor{red}{a}$

## Stochastic Policy

## Deterministic Policy

**Policy:**

$$\pi(\textcolor{red}{a}|\textcolor{green}{s}; \boldsymbol{\theta})$$

$$\pi(\textcolor{green}{s}; \boldsymbol{\theta})$$

**Output:**

Probability distribution  
over the action space

Action  $\textcolor{red}{a}$

**Control:**

Randomly sample an action  
from the distribution

Directly use  
the output,  $\textcolor{red}{a}$

**Application:**

Mostly discrete  
control

Continuous  
control