



# Markov Decision Process Examples

# Example I: Maze Problem as MDP




- Rewards:  $-1$  per time-step
- Actions: N, E, S, W
- States: Agent's location

# Maze Problem as MDP




- Numbers represent value  $V^\pi(s)$  of each state  $s$

# Maze Problem as MDP



- Arrows represent policy  $\pi(s)$  for each state  $s$

# Maze Problem as MDP

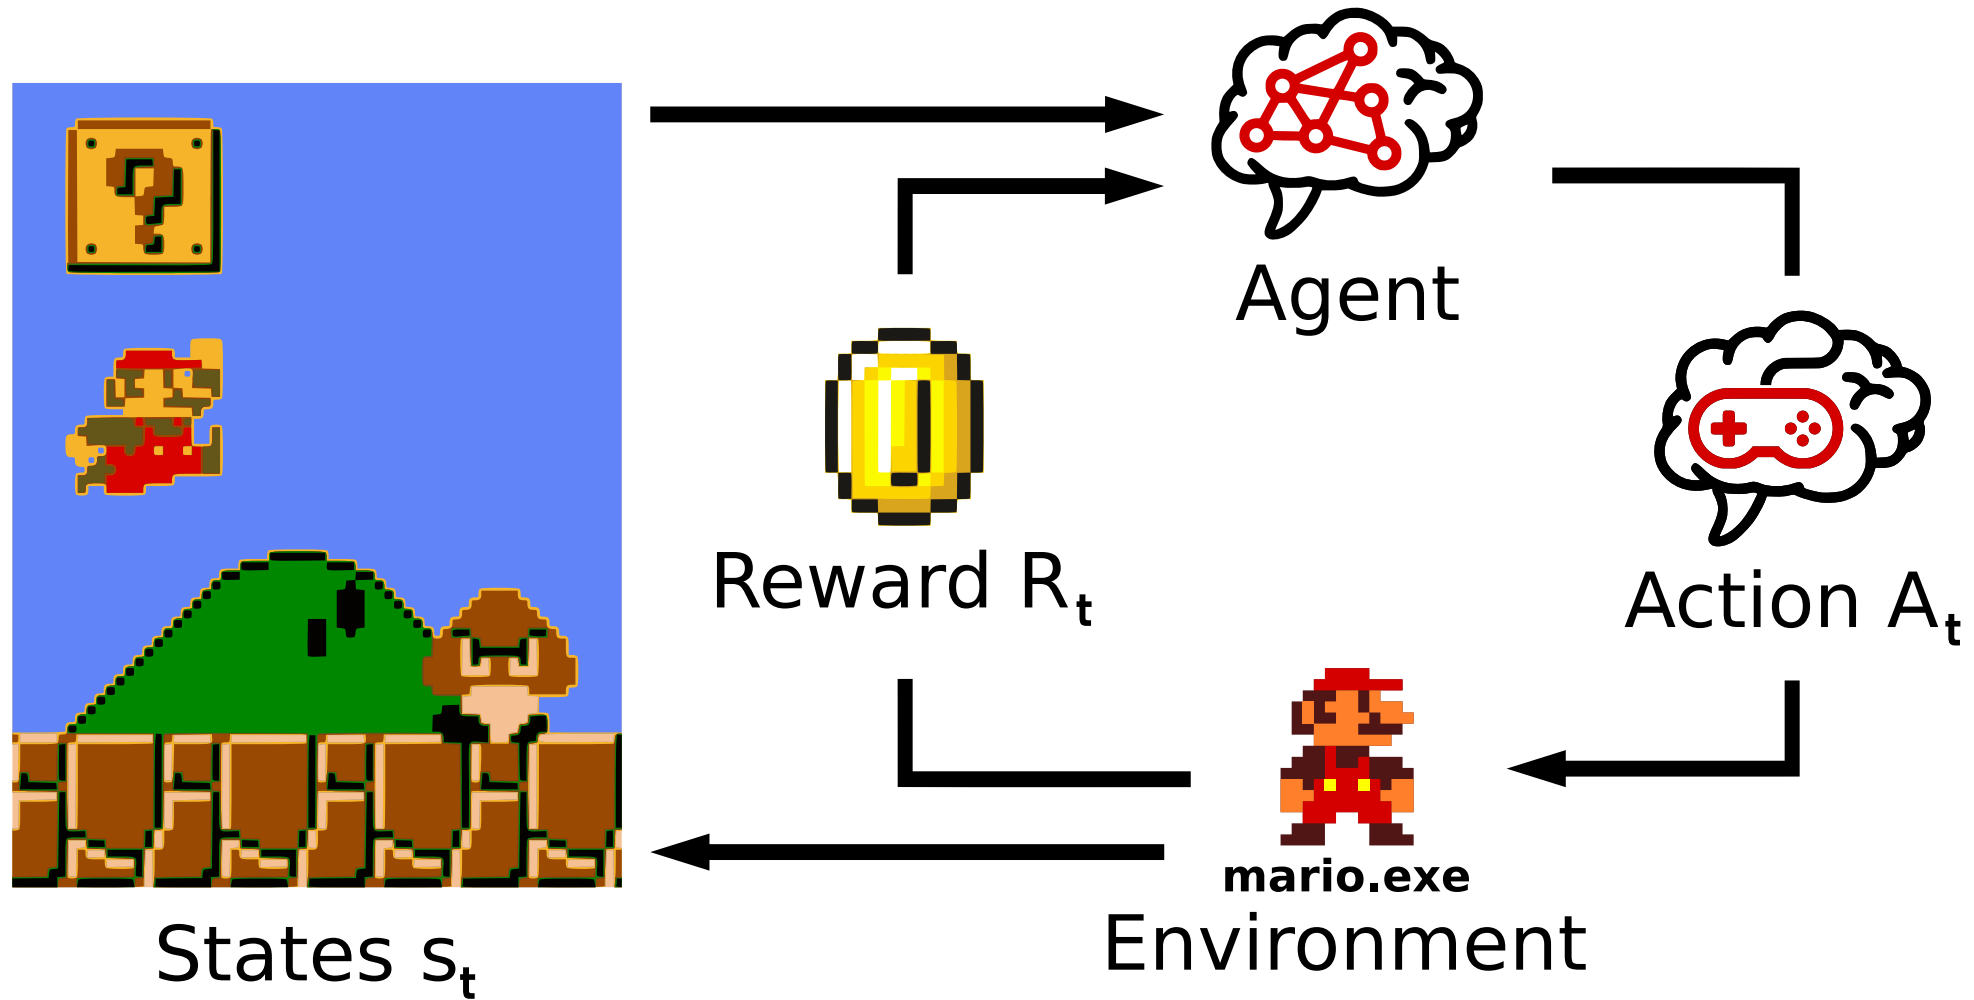


<https://www.samyzaf.com/ML/rl/qmaze.html>

# Example II: Super Mario Bros. Game




# Super Mario Bros. Game as MDP



# Terminology: **state** and **action**

state  $s$  (this frame)

Action  $a \in \{\text{left}, \text{right}, \text{up}\}$





# Terminology: state transition



## state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.

# Terminology: state transition

## state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.
- $p(s'|s, a) = \mathbb{P}(S' = s' | S = s, A = a)$ .





# Terminology: reward


## reward $R$



- Collect a coin:  $R = +1$
- Win the game:  $R = +10000$
- Touch a Goomba:  $R = -10000$  (game over).
- Nothing happens:  $R = 0$



# Terminology: policy

policy  $\pi$



- $\pi(a | s)$  is the probability of taking action  $A = a$  given state  $s$ , e.g.,
  - $\pi(\text{left} | s) = 0.2$ ,
  - $\pi(\text{right} | s) = 0.1$ ,
  - $\pi(\text{up} | s) = 0.7$ .
- Upon observing state  $S = s$ , the agent's action  $A$  can be random.


# Example III: Robot in Grid World






# Example: Robot in Grid World

- A maze-like problem
  - The agent lives in a grid
  - Walls block the agent's path
- Noisy movement: actions do not always go as planned
  - 80% of the time, the action North takes the agent North (if there is no wall there)
  - 10% of the time, North takes the agent West; 10% East
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
  - Small "living" reward each step (can be negative)
  - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards




# Grid World Actions


Deterministic Grid World



Stochastic Grid World




# Markov Decision Process



- **States:** Positions of Robots
- **Actions:** Movements
- **State Transitions:**  
Uncertain resulting state for the chosen direction
- **Rewards:** small living reward in non-terminate states; big rewards of +1 or -1 in terminate states




# Value Iteration



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=0$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=1



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=2



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=3



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=4$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=5



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=6



Noise = 0.2  
Discount = 0.9  
Living reward = 0



$k=7$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

k=8



Noise = 0.2  
Discount = 0.9  
Living reward = 0

k=9



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=10



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=11



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=12




Noise = 0.2  
Discount = 0.9  
Living reward = 0

# k=100




Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Policy Extraction



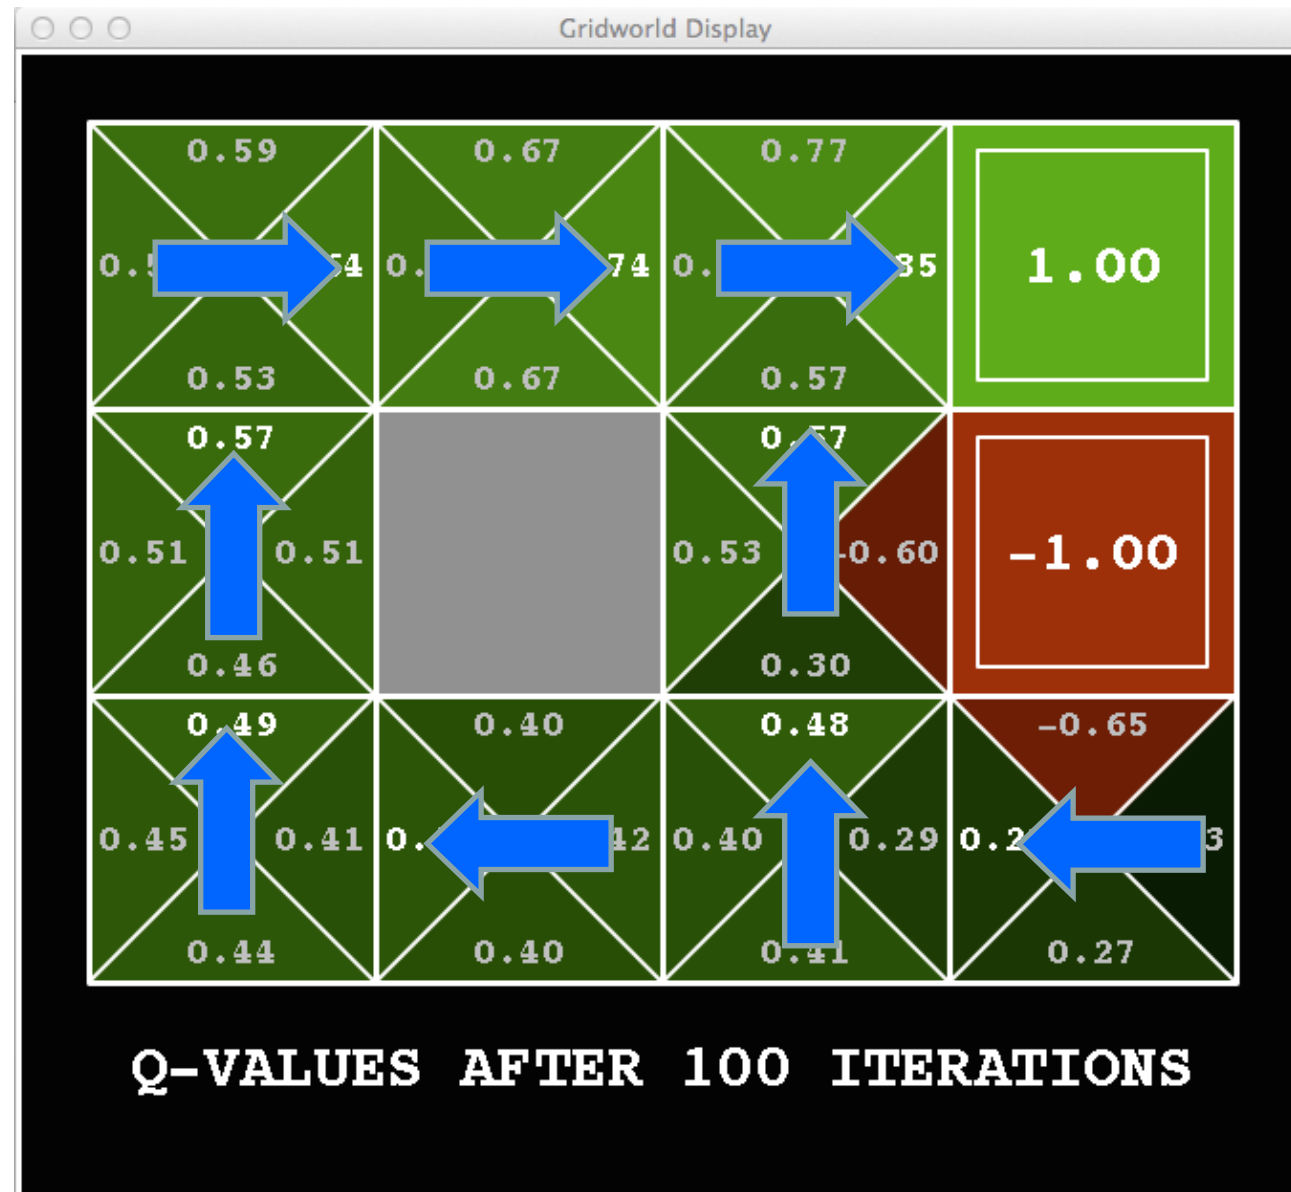


# Q Values



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Optimal Policy



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Example IV: Jack's Car Rental Problem



- States: Two locations, maximum of 20 cars at each
- Actions: Move up to 5 cars between locations overnight
- Reward: \$10 for each car rented (must be available)
- Transitions: Cars returned and requested randomly
  - Poisson distribution,  $n$  returns/requests with prob  $\frac{\lambda^n}{n!} e^{-\lambda}$
  - 1st location: average requests = 3, average returns = 3
  - 2nd location: average requests = 4, average returns = 2



# Shared-bike Relocation Problem



# Policy Iteration

